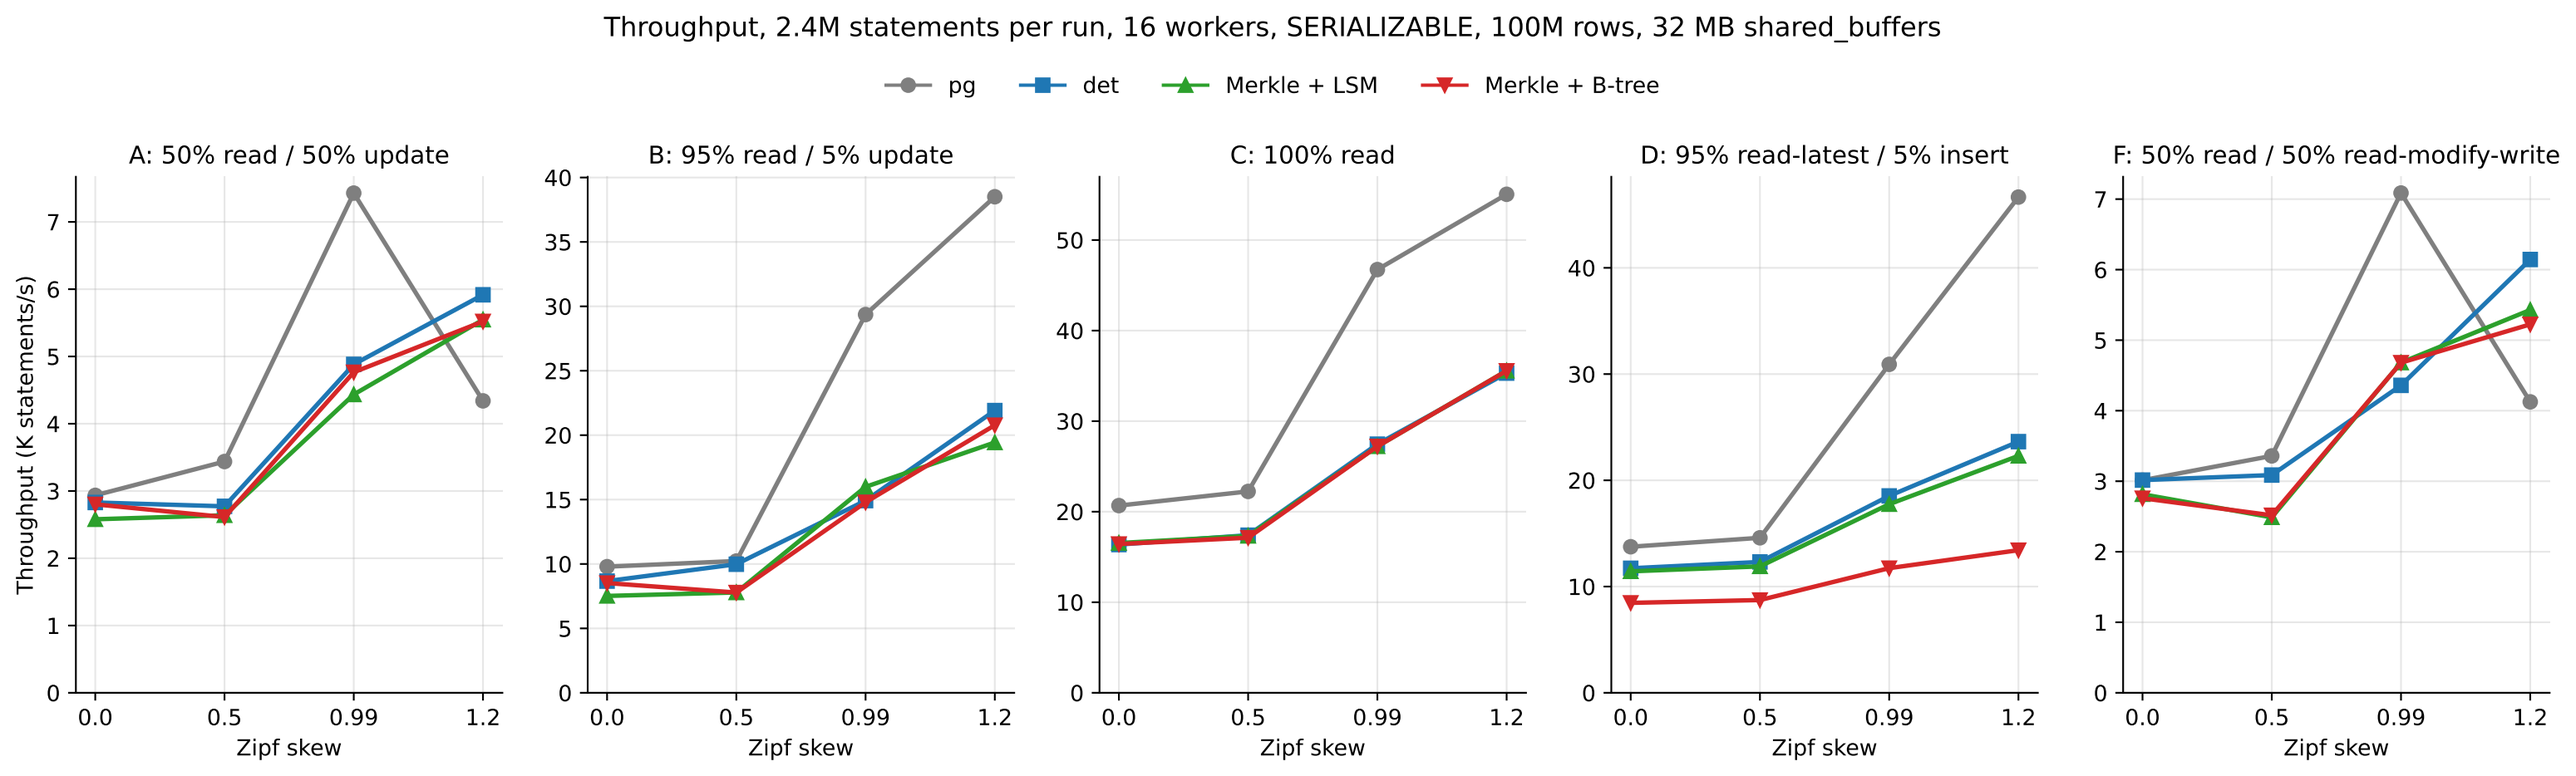
Throughput, 2.4M statements per run, 16 workers, SERIALIZABLE, 100M rows, 32 MB shared_buffers
pg
det
Merkle + LSM
Merkle + B-tree
A: 50% read / 50% update
B: 95% read / 5% update
C: 100% read
D: 95% read-latest / 5% insert
F: 50% read / 50% read-modify-write
Throughput (K statements/s)
Zipf skew
0.0
0.5
0.99
1.2