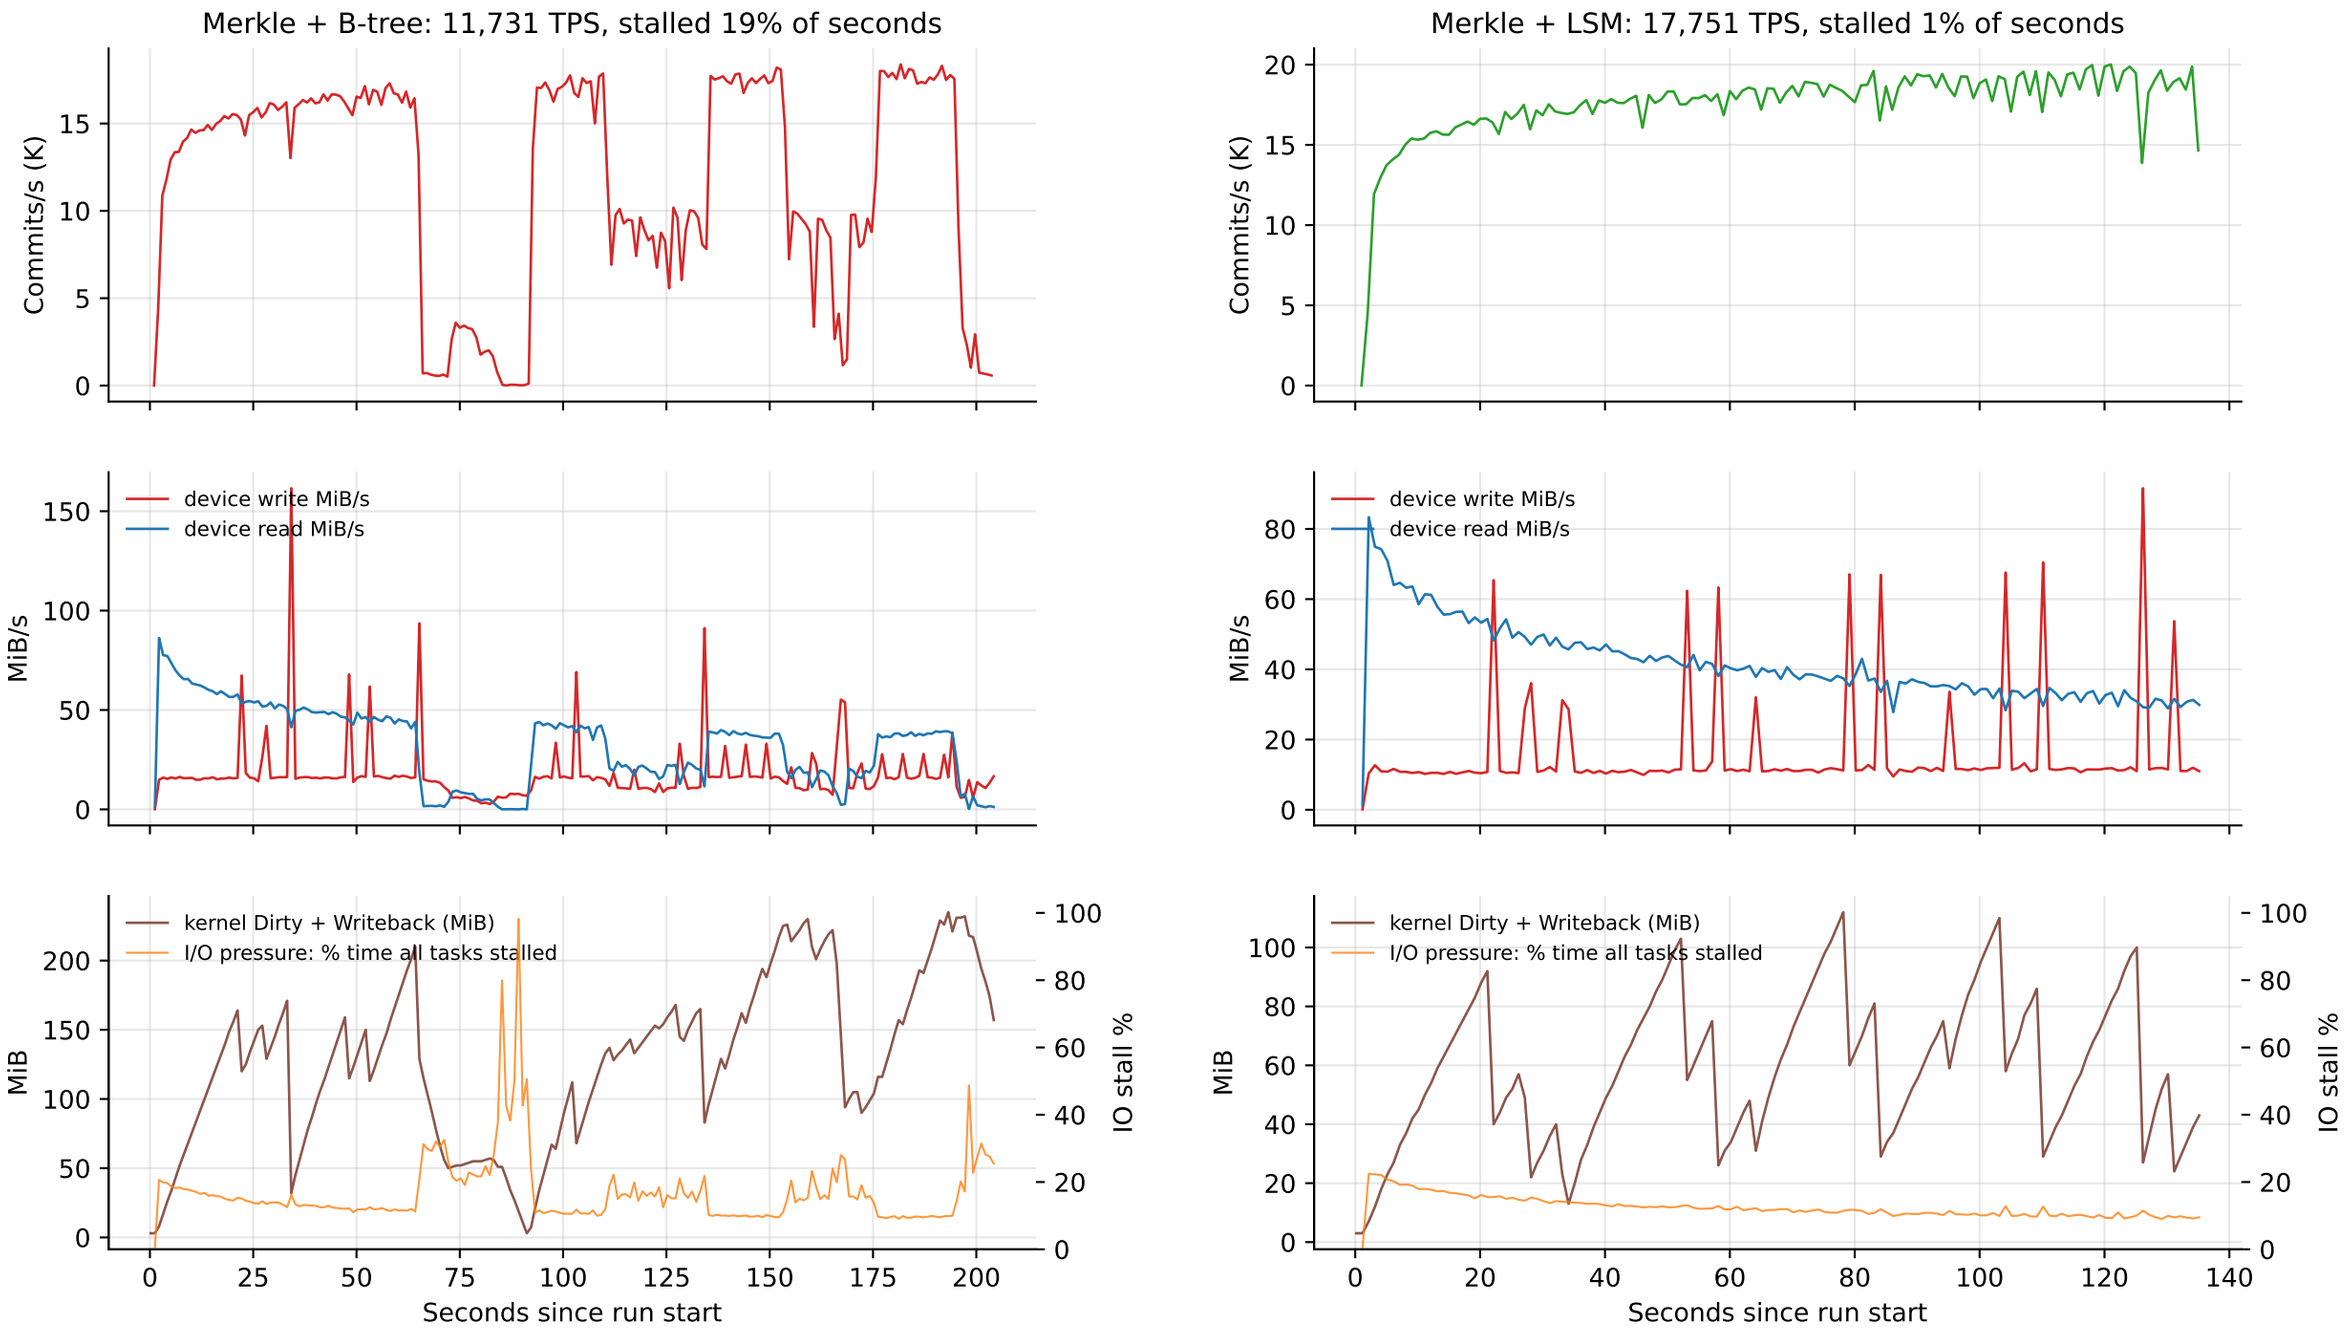Workload D, skew 0.99: during the B-tree stalls the device serves almost no reads or writes while I/O pressure peaks; kernel dirty pages stay small (<250 MiB), so this is the device stalling, not a writeback backlog. The LSM run (sequential L0 appends instead of random writes into the 3 GB hash-ordered index) never stalls. The SSD-internal cause is not observable here.